

Import required packages

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

Read the data

```
df=pd.read_csv("/content/drive/MyDrive/Colab Notebooks/travel
insurance claim.csv")
```

Performing EDA

```
df.shape
```

```
(50553, 12)
```

```
df.head()
```

	ID	Agency	Agency Type	...	Commision (in value)	Gender	Age
0	3433	CWT	Travel Agency	...	17.82	NaN	31
1	4339	EPX	Travel Agency	...	0.00	NaN	36
2	34590	CWT	Travel Agency	...	11.88	NaN	75
3	55816	EPX	Travel Agency	...	0.00	NaN	32
4	13816	EPX	Travel Agency	...	0.00	NaN	29

```
[5 rows x 12 columns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 50553 entries, 0 to 50552
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	50553 non-null	int64
1	Agency	50553 non-null	object
2	Agency Type	50553 non-null	object
3	Distribution Channel	50553 non-null	object
4	Product Name	50553 non-null	object
5	Claim	50553 non-null	int64
6	Duration	50553 non-null	int64
7	Destination	50553 non-null	object
8	Net Sales	50553 non-null	float64
9	Commision (in value)	50553 non-null	float64
10	Gender	14600 non-null	object
11	Age	50553 non-null	int64

```
dtypes: float64(2), int64(4), object(6)
memory usage: 4.6+ MB
```

Here as we can see that Gender Column Having more then 70% of NaN value So We Drop that Gender Column

```
df=df.drop(['Gender'],axis = 1)
df=df.drop(['ID'], axis=1)
```

```
df.describe()
```

	Claim	Duration	...	Commision (in value)
Age				
count	50553.000000	50553.000000	...	50553.000000
50553.000000				
mean	0.014658	49.425969	...	9.83809
40.011236				
std	0.120180	101.434647	...	19.91004
14.076566				
min	0.000000	-2.000000	...	0.00000
0.000000				
25%	0.000000	9.000000	...	0.00000
35.000000				
50%	0.000000	22.000000	...	0.00000
36.000000				
75%	0.000000	53.000000	...	11.55000
44.000000				
max	1.000000	4881.000000	...	283.50000
118.000000				

```
[8 rows x 5 columns]
```

```
df[df["Duration"] <0]
```

	Agency	Agency Type	...	Commision (in value)	Age
4063	JZI	Airlines	...	6.3	118
38935	JZI	Airlines	...	6.3	118
48367	JZI	Airlines	...	7.7	118

```
[3 rows x 10 columns]
```

```
df[df["Age"] > 100]
```

	Agency	Agency Type	...	Commision (in value)	Age
90	JWT	Airlines	...	31.20	118
108	JWT	Airlines	...	12.40	118
140	JWT	Airlines	...	15.60	118
153	JWT	Airlines	...	31.20	118
181	JWT	Airlines	...	12.40	118
...
50158	JWT	Airlines	...	24.00	118
50179	JWT	Airlines	...	12.40	118

50250	JWT	Airlines	...	12.40	118
50429	JZI	Airlines	...	12.25	118
50478	CCR	Travel Agency	...	9.57	118

[795 rows x 10 columns]

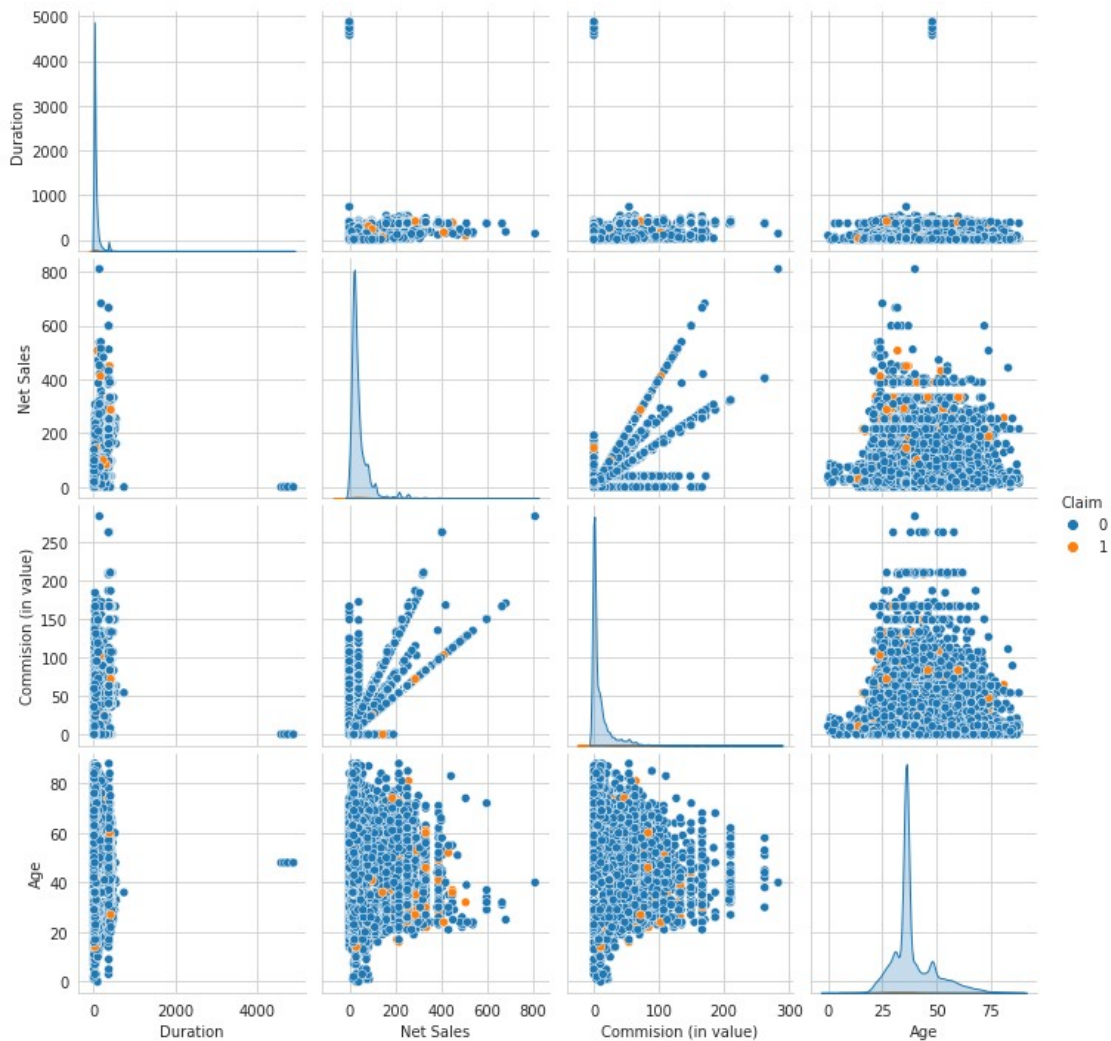
df[df["Net Sales"] < 0]

	Agency	Agency Type	...	Commision (in value)	Age
6	JZI	Airlines	...	24.15	26
128	EPX	Travel Agency	...	0.00	37
139	EPX	Travel Agency	...	0.00	46
173	CWT	Travel Agency	...	5.94	31
336	CWT	Travel Agency	...	11.88	31
...
50121	CWT	Travel Agency	...	59.40	45
50149	ART	Airlines	...	0.49	118
50177	CWT	Travel Agency	...	29.70	49
50394	JZI	Airlines	...	7.70	57
50399	CWT	Travel Agency	...	29.70	31

[528 rows x 10 columns]


```
df.loc[df['Duration'] < 0, 'Duration'] = 49.425969
df.loc[df['Age'] > 100, 'Age'] = 40.011236
df.loc[df['Net Sales'] < 0, 'Net Sales'] = 40.800977
```

```
sns.set_style("whitegrid");
sns.pairplot(df, hue="Claim");
plt.show()
```



```
print(df["Claim"].value_counts())
print("-----")
plt.figure(figsize=(7,7))
df["Claim"].value_counts().plot.pie(autopct="%.1f%%")
plt.show()

0    49812
1     741
Name: Claim, dtype: int64
-----
```



Looking For Correlations

```
ax = plt.subplots(figsize=(15, 10))
sns.heatmap(df.corr(), square=True, annot=True)
<matplotlib.axes._subplots.AxesSubplot at 0x7ffa3f8ccfd0>
```



Commision (in value) column having High Correlation between Feature 'Net Sales' and very Low Correlation with the 'Target' so we are Dropping The Commision (in value) Column

```
df=df.drop(['Commision (in value)'],axis=1)
```

```
df.head()
```

```

Agency  Agency Type  ... Net Sales  Age
0    CWT  Travel Agency  ...      0.0  31.0
1    EPX  Travel Agency  ...     69.0  36.0
2    CWT  Travel Agency  ...     19.8  75.0
3    EPX  Travel Agency  ...     20.0  32.0
4    EPX  Travel Agency  ...     15.0  29.0

```

```
[5 rows x 9 columns]
```

Separating categorical and numerical data

```
df_num = df.select_dtypes(["int64","float64"])
df_cat = df.select_dtypes("object")
```

For df_num

```
df_num.head()
```

	Claim	Duration	Net Sales	Age
0	0	7.0	0.0	31.0
1	0	85.0	69.0	36.0
2	0	11.0	19.8	75.0
3	0	16.0	20.0	32.0
4	0	10.0	15.0	29.0

```
df_cat.head()
```

	Agency	Agency Type	...	Product Name
0	CWT	Travel Agency	...	Rental Vehicle Excess Insurance
1	EPX	Travel Agency	...	Cancellation Plan
2	CWT	Travel Agency	...	Rental Vehicle Excess Insurance
3	EPX	Travel Agency	...	2 way Comprehensive Plan
4	EPX	Travel Agency	...	Cancellation Plan KOREA,

```
[5 rows x 5 columns]
```

```
df_num.describe()
```


	Claim	Duration	Net Sales	Age
count	50553.000000	50553.000000	50553.000000	50553.000000
mean	0.014658	49.428981	41.852794	38.784779
std	0.120180	101.433893	47.536249	10.049564
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	9.000000	18.500000	35.000000
50%	0.000000	22.000000	27.000000	36.000000
75%	0.000000	53.000000	48.000000	42.000000
max	1.000000	4881.000000	810.000000	88.000000

Skewness


```
from scipy.stats import skew
```

```
for col in df_num:
    print(col, ":-", skew(df_num[col]))
    plt.figure(figsize=(8,8))
    sns.distplot(df_num[col])
    plt.show()
```


Claim :- 8.076976414369875




Duration :- 22.872492167935484



Net Sales :- 3.751192202882967




Age :- 1.2031538221807883




```
#removing skewness
for col in df_num:
    if skew(df_num[col]) >= 0.5 or skew(df_num[col]) <= -0.5:
        df_num[col] = np.sqrt(df_num[col])

for col in df_num:
    print(col, ":-", skew(df_num[col]))
    plt.figure(figsize=(8,8))
    sns.distplot(df_num[col])
    plt.show()
```


Claim :- 8.076976414369875




Duration :- 2.409425954993082



Net Sales :- 1.4439989184434878





Age :- 0.6848711397302739





Bivariate Analysis against the target

```
for col in df_num:
    plt.figure(figsize=(8,8))
    sns.boxplot(data=df,x="Claim",y=col , hue='Claim')
    plt.show()
```









Scaling

```
from sklearn.preprocessing import MinMaxScaler

for col in df_num:
    ms = MinMaxScaler()
    df_num[col] = ms.fit_transform(df_num[[col]])

df_num.head()
```

	Claim	Duration	Net Sales	Age
0	0.0	0.037870	0.000000	0.593526
1	0.0	0.131964	0.291865	0.639602
2	0.0	0.047472	0.156347	0.923186
3	0.0	0.057254	0.157135	0.603023
4	0.0	0.045263	0.136083	0.574060

For df_cat

```
df_cat.head()
```

	Agency	Agency Type	...	Product Name
0	CWT	Travel Agency	...	Rental Vehicle Excess Insurance
1	EPX	Travel Agency	...	Cancellation Plan
2	CWT	Travel Agency	...	Rental Vehicle Excess Insurance
3	EPX	Travel Agency	...	2 way Comprehensive Plan
4	EPX	Travel Agency	...	Cancellation Plan KOREA,




Destination
MALAYSIA
SINGAPORE
MALAYSIA
INDONESIA
REPUBLIC OF

```
[5 rows x 5 columns]
```

```
plt.figure(figsize=(15,10))
plt.subplot(2,2,1)
sns.countplot(df['Agency Type'])
plt.title('Agency Type')
plt.subplot(2,2,2)
sns.countplot(df['Distribution Channel'])
plt.title('Distribution Channel')
```

```
plt.subplot(2,2,3)
sns.countplot(df['Agency'])
plt.xticks(rotation=90)
plt.title('Agency')
```

```
Text(0.5, 1.0, 'Agency')
```



```
plt.figure(figsize=(15,10))
```

```
wedges, texts = plt.pie(df['Destination'].value_counts(),
                        labels = df['Destination'].unique(),
```


```
                        shadow = True,
```

```
                        textprops = dict(color = "magenta"))
```

```
plt.pie(df['Destination'].value_counts(), labels=df['Destination'].unique())
```

```
plt.title('Destination')
```

```
Text(0.5, 1.0, 'Destination')
```



```
plt.figure(figsize=(15,10))
```

```
wedges, texts = plt.pie(df['Product Name'].value_counts(),
                        labels = df['Product
```

```
Name'].unique(),
```

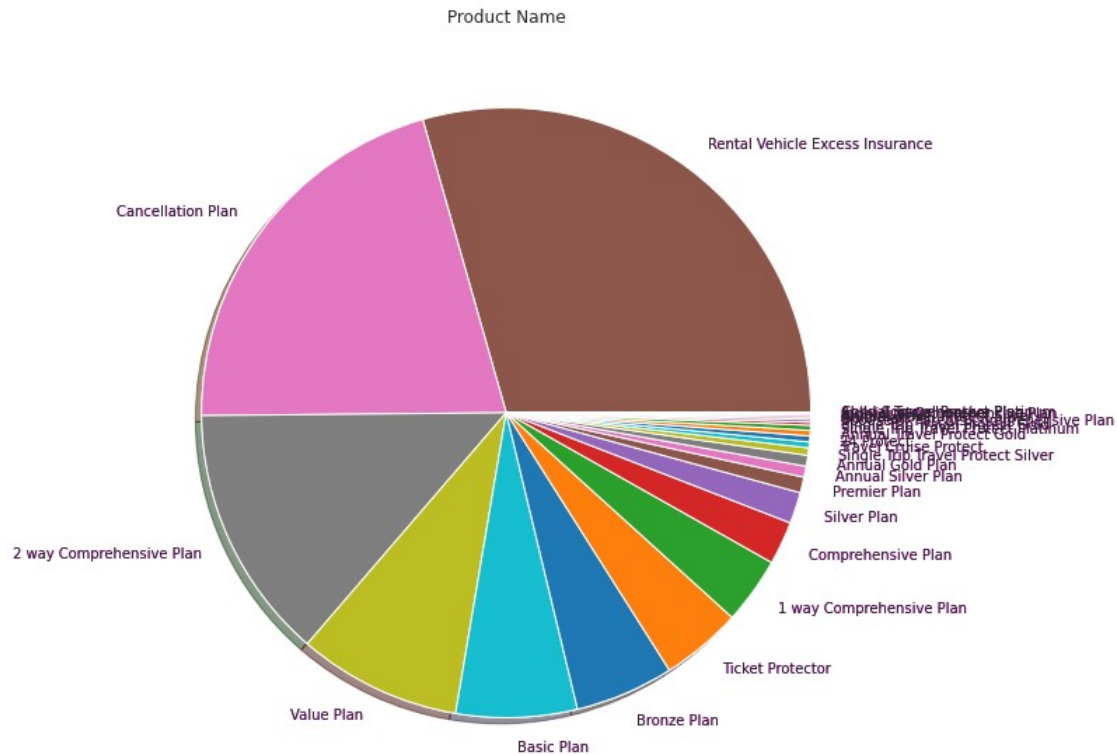
```
shadow = True,
```

```
textprops = dict(color = "magenta"))
```

```
plt.pie(df['Product Name'].value_counts(), labels=df['Product
Name'].unique())
```

```
plt.title('Product Name')
```

```
Text(0.5, 1.0, 'Product Name')
```



Encoding

```
from sklearn.preprocessing import LabelEncoder

for col in df_cat:
    le = LabelEncoder()
    df_cat[col] = le.fit_transform(df_cat[[col]])

df_cat.head()
```

	Agency Destination	Agency Type	Distribution Channel	Product Name
0	6	1	1	16
56				
1	7	1	1	10
79				
2	6	1	1	16
56				
3	7	1	1	1
38				
4	7	1	1	10
47				

Concatenate Both df_num and df_cat

```
df_new = pd.concat([df_num,df_cat], axis=1)
```

```
df_new.head()
```

```
      Claim  Duration  Net Sales  ...  Distribution Channel  Product Name
Destination
0      0.0   0.037870   0.000000  ...                    1             16
56
1      0.0   0.131964   0.291865  ...                    1             10
79
2      0.0   0.047472   0.156347  ...                    1             16
56
3      0.0   0.057254   0.157135  ...                    1              1
38
4      0.0   0.045263   0.136083  ...                    1             10
47
```

```
[5 rows x 9 columns]
```

```
X = df_new.drop("Claim",axis=1)
y = df_new["Claim"]
```

Train-Test Splitting bold text

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.3,random_state=1)
```

#Baseline Model

LogisticRegression

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression()
```

```
model.fit(X_train,y_train)
```

```
LogisticRegression()
```

```
y_pred = model.predict(X_test)
```


```
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, classification_report, confusion_matrix
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	14961
1.0	0.00	0.00	0.00	205
accuracy			0.99	15166
macro avg	0.49	0.50	0.50	15166

weighted avg	0.97	0.99	0.98	15166
--------------	------	------	------	-------

```
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
plt.show()
```



Using RandomOverSampler to Balanced the training data

```
from imblearn.over_sampling import RandomOverSampler
ros = RandomOverSampler(random_state=1)
X_sample1, y_sample1 = ros.fit_resample(X_train,y_train)
pd.Series(y_sample1).value_counts()

1.0    34851
0.0    34851
Name: Claim, dtype: int64

lr2 = LogisticRegression()
lr2.fit(X_sample1, y_sample1)


LogisticRegression()

y_pred1= lr2.predict(X_test)
print(classification_report(y_test,y_pred1))
```

	precision	recall	f1-score	support
0.0	0.99	0.82	0.90	14961

	1.0	0.05	0.63	0.08	205
accuracy				0.81	15166
macro avg		0.52	0.73	0.49	15166
weighted avg		0.98	0.81	0.89	15166

```
sns.heatmap(confusion_matrix(y_test, y_pred1), annot=True, fmt='d')
plt.show()
```



under sampling

```
from imblearn.under_sampling import RandomUnderSampler
rus = RandomUnderSampler(random_state=1)
X_sample2, y_sample2 = rus.fit_resample(X_train,y_train)
pd.Series(y_sample2).value_counts()

1.0    536
0.0    536
Name: Claim, dtype: int64


lr3 = LogisticRegression()
lr3.fit(X_sample2, y_sample2)

LogisticRegression()

y_pred2 = lr2.predict(X_test)
print(classification_report(y_test,y_pred2))
```

	precision	recall	f1-score	support
0.0	0.99	0.82	0.90	14961
1.0	0.05	0.63	0.08	205
accuracy			0.81	15166
macro avg	0.52	0.73	0.49	15166
weighted avg	0.98	0.81	0.89	15166

```
from sklearn.metrics import confusion_matrix, classification_report
sns.heatmap(confusion_matrix(y_test, y_pred2), annot=True, fmt='d')
plt.show()
```



#Decision tree Here we use DecisionTree with RandomOverSampler Technique to Balanced The DataFrame and to get classification report

```
from sklearn.tree import DecisionTreeClassifier
ros = RandomOverSampler(random_state=1)
X_sample3, y_sample3 = ros.fit_resample(X_train,y_train)
pd.Series(y_sample3).value_counts()
1.0    34851
0.0    34851
Name: Claim, dtype: int64

dtc=
DecisionTreeClassifier(max_depth=8 ,min_samples_leaf=50,criterion="ent
```

```

copy")
dtc.fit(X_sample3, y_sample3)

DecisionTreeClassifier(criterion='entropy', max_depth=8,
min_samples_leaf=50)

y_pred3 = dtc.predict(X_test)
print(classification_report(y_test,y_pred3))

```

	precision	recall	f1-score	support
0.0	0.99	0.77	0.87	14961
1.0	0.04	0.62	0.07	205
accuracy			0.77	15166
macro avg	0.51	0.70	0.47	15166
weighted avg	0.98	0.77	0.86	15166

```

sns.heatmap(confusion_matrix(y_test, y_pred3), annot=True, fmt='d')
plt.show()

```



Random forest

```

from sklearn.ensemble import RandomForestClassifier
ros = RandomOverSampler(random_state=1)
X_sample4, y_sample4 = ros.fit_resample(X_train,y_train)
pd.Series(y_sample4).value_counts()

```

```

1.0    34851
0.0    34851
Name: Claim, dtype: int64

rtc=
RandomForestClassifier(n_estimators=100,random_state=1,max_features=8,
max_depth=5)
rtc.fit(X_sample4, y_sample4)

RandomForestClassifier(max_depth=5, max_features=8, random_state=1)

y_pred4 = rtc.predict(X_test)
print(classification_report(y_test,y_pred4))


              precision    recall  f1-score   support

    0.0         0.99      0.82      0.90      14961
    1.0         0.05      0.63      0.09         205

 accuracy          0.82      15166
 macro avg         0.52      0.72      0.49      15166
 weighted avg         0.98      0.82      0.89      15166

sns.heatmap(confusion_matrix(y_test, y_pred4), annot=True, fmt='d')
plt.show()

```



#Gradient Boosting Classifier

```
#Gradient Boost
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
ros = RandomOverSampler(random_state=1)
```

```
X_sample5, y_sample5 = ros.fit_resample(X_train,y_train)
```

```
pd.Series(y_sample5).value_counts()
```

```
1.0    34851
```

```
0.0    34851
```

```
Name: Claim, dtype: int64
```

```
gb=GradientBoostingClassifier(n_estimators=100)
```

```
gb.fit(X_sample5, y_sample5)
```


```
GradientBoostingClassifier()
```

```
y_pred5 = gb.predict(X_test)
```

```
print(classification_report(y_test,y_pred5))
```

	precision	recall	f1-score	support
0.0	0.99	0.82	0.90	14961
1.0	0.05	0.61	0.08	205
accuracy			0.82	15166
macro avg	0.52	0.72	0.49	15166
weighted avg	0.98	0.82	0.89	15166

```
sns.heatmap(confusion_matrix(y_test, y_pred5), annot=True, fmt='d')  
plt.show()
```



#Adaboost Boosting Classifier

```
from sklearn.ensemble import AdaBoostClassifier
ros = RandomOverSampler(random_state=1)
X_sample6, y_sample6 = ros.fit_resample(X_train,y_train)
pd.Series(y_sample5).value_counts()
1.0    34851
0.0    34851
Name: Claim, dtype: int64
ab=AdaBoostClassifier(n_estimators=100)
ab.fit(X_sample6, y_sample6)
AdaBoostClassifier(n_estimators=100)
y_pred6 = ab.predict(X_test)
print(classification_report(y_test,y_pred6))
```

	precision	recall	f1-score	support
0.0	0.99	0.80	0.89	14961
1.0	0.04	0.66	0.08	205
accuracy			0.80	15166
macro avg	0.52	0.73	0.48	15166

weighted avg 0.98 0.80 0.88 15166

```
sns.heatmap(confusion_matrix(y_test, y_pred6), annot=True, fmt='d')  
plt.show()
```

